

Homework #2: Virtual World at Facecook

Due on Thursday, September 19 at 11:59 p.m.

The Facecook CEO has come up with a new concept for the Facecook website: applications that allow Facecook users to interact in rich, virtual environments. Since this concept was articulated, the Facecook framework design committee has developed a simulation framework for simple 2D worlds that evolve in real time. After your successful work on a social network graph last week, your manager promoted you to enhance a prototype virtual environment that your co-workers have been working on.

The prototype environment is designed to simulate nature. Grass will grow in the warm sunshine of a world, where rabbits hop playfully. Of course, no social world would be complete without conflict; thus, there will be foxes constantly plotting to catch a rabbit for dinner.

Your goals for this assignment are:

- Effectively use objects and inheritance for encapsulation and code reuse
- Program implementations for existing interfaces
- Experience working with a medium-sized application
- Start thinking about code from a design point-of-view

Instructions

Luckily, your co-workers have already developed a virtual world environment to be used. Your job is to build the objects that will fill the world. You will need to build rabbits and foxes, as well as the AI for these objects. Furthermore, you must create an animal of your choice, whose behavior you must define.

Documentation

The Facecook design team has provided the following documentation for you. Please read this carefully.

(1) Types of Objects in the World – Specifications and Definitions

The world contains the following object types, with each type having different properties and specifications described here.

Actor: An **Actor** is a participant in the world. **Actors** are objects which are able to act in the world, and they have a cool down property and a view range.

An **Actor**'s view range is defined as the number of spaces away from the actor that is visible. This means that if an **Actor** is currently at a location (i, j) in our world and has a view range of k (k is a nonnegative integer), then the **Actor** can see all spaces in the square defined by the corners $(i - k, j - k)$, $(i + k, j - k)$, $(i + k, j + k)$, $(i - k, j + k)$, provided that all of the corners of that square is within the boundary of the world (if not, then set of spaces that the **Actor** can see is the largest subset of spaces of that square such that the subset is within the boundary of the world). Note that we define $(0, 0)$ to be the top left corner of the world.

Time in the virtual world progresses in discrete time steps and actors also have a cool down property, an integer that determines how often an actor acts. The actor's cooldown determines how many steps must pass before the actor can act again. For example, if an actor has a cooldown of 0 and it acts on step 0, it may act again on step 1. If an actor had a cooldown of 1 and it acts on step 0, however, then it may not act again until step 2. The world implements these cooldown properties for you; you only need to understand them so you know how long an actor must wait before acting again, to help you build a better AI.

Animal: An **Animal** is an **Actor** that has an added property of energy, and is able to eat, breed, and move. Energy is used to act in the world; every time an **Animal** acts, it loses one unit of energy. More specifically, any time an **Animal** uses the AI to select an action (please read the later section on AIs), the **Animal** loses one unit of energy, even if the AI returns a **null** response. If an **Animal**'s energy reaches zero, the **Animal** dies and it removes itself from the world. Each kind of **Animal** will have a different initial level of energy, and a different maximum energy. The latter is the maximum energy permitted at the end of a step, after considering the energy used to ask the AI what to do, and considering any energy gained by eating.

Edible: An **Edible** object is an object which an **Animal** can eat. **Edible** objects have an energy value that specify how much energy the eater gains when the object is eaten. For example, suppose grass has an energy value of 5. Each time a **Rabbit** eats grass it would then gain 5 units of energy. An **Edible**'s energy value should be constant. Note that the energy value of an **Edible** is different from (though related to) the current energy and maximum energy of an **Animal**.

(2) Specific Objects in the World

You will create some of the following classes that implement the interfaces described above. Many of the following have been created for you.

Grass: **Grass** is an **Edible Actor**. The **Grass** class has been implemented for you.

Gardener: A **Gardener** is a simple **Actor** which stays in place and adds **Grass** to the world at random locations. The **Gardener** class has been implemented for you.

Gnat: A **Gnat** is an **Actor**. **Gnats** do not lose energy when they act and therefore live forever. The **Gnat** class is provided only as sample code to demonstrate how **Actors** can be implemented.

Rabbit: A `Rabbit` is an `Animal` which is also `Edible`. Furthermore, `Rabbits` need to breed as often as possible and only eat `Grass`. You must write your own `RabbitImpl` class which implements the `Rabbit` interface.

Fox: A `Fox` is an `Animal` which eats `Rabbits`. `Foxs` also breeds at a lower rate than `Rabbits`. You must write your own `FoxImpl` class which implements the `Fox` interface.

Your Third Creature: The Facebook team is giving you full creative license to design and implement a third creature of your choice. Your creature does not need to be an animal but it should be interesting—more interesting than grass! It should be an actor that has an AI, like the `Fox` and the `Rabbit`. Feel free to have fun with this one; for example, you could implement an alien creature that randomly teleports around the world.

We have also provided some constants for you which we think will make a stable world. In our opinion, reasonably smart `Rabbits` and `Foxes` inhabiting this world will allow it to run indefinitely without either species going extinct. You may change the existing constant values to make your world more stable.

(3) Commands and Behavior – Communication

The AI uses `Command` in order to tell `Actors` what to do. `Animals` have three types of `Commands`: `BreedCommand`, `EatCommand`, and `MoveCommand`, each of which call the `Animal`'s `breed`, `eat`, and `move` methods respectively:

Breeding: When an `Animal` breeds, it makes a copy of itself on an adjacent tile. The `Animal` can only breed when it has enough energy to do so. Breeding occurs alone; there is no mating between multiple `Animal`. When it breeds, an `Animal`'s energy is reduced to 50% of its former energy (rounded down), and the newly-bred `Animal` also starts at 50% of its parent's energy. Finally, newly bred `Animal` must be placed in an empty location that is adjacent to the parent.

Eating: An `Animal` is only able to eat something that is adjacent to it. When eating an item, the `Animal` removes the item from the world and remains in place. The eaten item must be `Edible`, and the `Animal` gains its energy value.

Moving: An `Animal` can only move one space at a time and movement is restricted to only adjacent spaces. An `Animal` must move only to empty spaces.

Dying: Dying is also something you must implement. An `Animal` dies when its energy reaches 0 and you must remove the `Animal` from the world when it reaches that energy level. No dead `Animals` should be left rotting in the world and an `Actor` is expected to remove itself when it dies; it should remove itself if, after its last action, it has 0 energy. This means that if an `Animal` had 1 energy and used that last energy to eat something (and successfully gained energy from the `Edible`), the `Animal` would not die because it would have more than 0 energy at the end of the action.

(4) The AI – Decision Making

How you implement your AI determines how **Rabbits** and **Foxs** behave in the world.

We describe how the AI works using the **Gnat** code as an example. When a **Gnat** is instantiated, it stores a **GnatAI** object in a data field. When a **Gnat** acts, its **act** method returns a **Command** object (or **null**), with which the AI will use to execute the gnat's next action.

Your AI implementations must implement the AI interface and have a zero-argument constructor (or no constructor, in which case Java implicitly defines a zero-argument constructor for you). Our code that evaluates your submission requires that all AIs have zero-argument constructors, and you can not participate in our virtual world tournament (see below) if your AIs require constructor arguments.

The AI will determine what command to provide based on the **Actor** and the state of the world; the visible portion of the world depends on the **Actor**. For example, if a **Rabbit** has a view range of 1 and is looking for **Grass**, it will only be able to see locations that are within 1 adjacent square (including diagonals). The AI should be able to see all these locations and then determine which **Command** it should create depending on the situation.

If the AI cannot determine a proper command to execute, it may return a **null** value, which commands the **Actor** to do nothing. All actors lose energy each time **act** is called, even if the actor chooses to do nothing.

(5) The World—Putting it all together

The **World** is an $n \times n$ grid of **Actors**, with the top left corner being (0,0). Each space in the grid can hold a single object. Empty spaces on the grid are represented by a **null** value and each coordinate in the **World** is represented by a **Location** object. **Location** objects contain some helpful methods you should be aware of. Furthermore the **World** provides the notion of direction that will aide interaction with the **Location** object.

The **World** has several methods. Given a **World** object, anything is able to add, remove, search, and view objects in the world. Furthermore, the **World** provides a step function that you will not use personally but should understand how it works. Each step forces every **Actor** in the **World** to act once, provided their cooldown allows them to act.

We have also provided a GUI you may use to visualize the world and see your **Animal**. The GUI has a simple interface containing two buttons: a “Step” to execute a single step; and a “Start/Stop” toggle button to run indefinitely until the toggle button is pressed again. For completing the homework, it is not necessary to understand the implementation of the GUI. You may run our code by running the main method in the **WorldUI** class.

(6) Customizing the World

We have provided an initial version of a class named **WorldLoader**. This class defines how the world should be set up, adding the initial elements into the world such as the starting number of rabbits, foxes, grass, initial energy of creatures, etc. You get to decide how many of each object are added. Choices here are not extremely important; it's just a way for you to see test your AIs in different situations.

For this homework, we understand that testing code in this virtual environment is very difficult, and we rather you devote your time to writing good AIs and understanding inheritance. You should try to run your implementations in the GUI, you can experiment with different values for the constants, and you write can write test code. However, **we do not have any testing-related requirements**. It is possible to earn full credit for this homework without writing any tests for your implementation.

Evaluation

Overall this homework is worth 100 points, plus up to 10 points of extra credit. To earn full credit you must do the following:

- Design your **Rabbit** and **Fox** to reuse as much code as possible. Avoid copying-and-pasting code; instead, design your code with useful abstractions and class hierarchies.
- Your AIs must implement the provided AI interface and must be independent of the rest of your implementation. The AI may rely on the **Fox** and **Rabbit** interfaces we provide with the assignment, but they should not depend on your specific implementation of those interfaces. For the arena-style tournament, your AI will be used to control our own fox and rabbit implementations. If you unsure whether your implementation will work in our arena, please ask the course staff using a private question on Piazza.
- You must name your Fox AI class **FoxAI** and your Rabbit AI class **RabbitAI**, including the exact capitalization given here. Both classes must either have no constructor or a public constructor without parameters. Your AI implementations must be in the `edu.cmu.cs.cs214.hw2.ai` package.
- You must name your Fox class **FoxImpl** and your Rabbit class **RabbitImpl**. They must be in the `edu.cmu.cs.cs214.hw2.actors` package.
- In general, adhere to the code organization. You may add classes, abstract classes, and interfaces that you desire to the code base. Place any new files in the appropriate package. For example, interfaces go in **interfaces**, **FoxImpl** goes in **actors**, etc.
- You must not delete, add, or modify any files in the **staff** directories.
- Use the most restrictive access level that makes sense for each of your fields and methods (i.e. use **private** unless you have a good reason not to). Instead of manipulating class fields directly, make them **private** and implement getter and setter methods to manipulate them from outside of the class. See [Controlling Access to Members of a Class](#) for a reference.
- As in previous assignments, provide reasonable documentation. Document at least each public class and public method with a short javadoc comment.
- Do not use magic numbers in your code. Magic numbers are unnamed constant values that do not have a well-defined meaning. For example, in:

```
for (int i = 0; i < 4; i++) {
    Direction d = Direction.getValues()[i];
    System.out.println(d);
}
```

In the above example, the number 4 is a magic number. Magic numbers make code hard to understand and add complexity when changing your code. Instead of using magic numbers, define `static final` constant variables with good names, such as `NUM.DIRECTIONS` in this example.

- As usual, make sure your code is readable. Use proper indentation and whitespace, abide by standard Java naming conventions, and add additional comments as necessary to document your code. Hint: using **Ctrl + Shift + F** to auto-format your code!

Additional hints:

- The tasks may be underspecified. In case of doubt use your judgment. If you want to communicate your assumptions, use comments in the source code.
- Do not use the `java.lang.Class` class or the `java.lang.reflect` package. You do not need – and should not use – those advanced techniques for this homework.

We will use the following approximate rubric to grade your work:

- Correct working `Rabbit` and `Fox` implementations with working AIs: 50 points
- Completion of your third animal implementation: 10 points
- Good design, including well designed class hierarchies and avoiding code replication: 30 points
- Documentation and style: 10 points
- Bonus: up to 10 points. After the homework deadline, we will run all students' `Rabbit` and `Fox` AIs in our own virtual arena with our own `Fox` and `Rabbit` implementations and award bonus points to the best competitors. Our animal implementations may use different values for the cooldown periods, the breed limit and similar values. Our `World` implementation will strictly enforce the `World` rules such as the `Animal` view limits. Your solution will be disqualified from the tournament if your implementations do not respect the rules (such as the view limit) of our `World`.

Have fun!